****
Lecturas 31/08: LLMs
****
Language Models are Few-Shot Learners
Tom B. Brown et al.
https://papers.nips.cc/paper/2020/file/1457c0d6bfcb4967418bfb8ac142f64a-Paper.pdf

Attention Is All You Need
Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, Illia Polosukhin
https://arxiv.org/abs/1706.03762
Vicente Vásquez

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding
Jacob Devlin, Ming-Wei Chang, Kenton Lee, Kristina Toutanova
https://arxiv.org/abs/1810.04805
Rodrigo Ugarte

****
Lecturas 02/9: Post-training
****
Training language models to follow instructions with human feedback
Long Ouyang et al.
https://arxiv.org/abs/2203.02155

Direct preference optimization: Your language model is secretly a reward model
Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D Manning, Chelsea Finn
https://arxiv.org/abs/2305.18290
Nicolas Lindl

LIMA: Less Is More for Alignment
Zhou et al.
https://arxiv.org/abs/2305.11206
Diego Toledo

****
Lecturas 07/9: Data
****
The Pile: An 800GB Dataset of Diverse Text for Language Modeling
Leo Gao et al.
https://arxiv.org/abs/2101.00027

Tülu 3: Pushing Frontiers in Open Language Model Post-Training (hasta pag 20.)
Lambert et al.
https://arxiv.org/pdf/2411.15124
Fernando Fajardo

Textbooks Are All You Need II: phi-1.5 technical report
Li et al.
https://arxiv.org/abs/2309.05463
María Dubo

****
Lecturas 09/9: Self-attention

****
In-Context Learning and Induction Heads
Olsson et al.
https://arxiv.org/abs/2209.11895

Mistral 7B
Albert Q. Jiang et al.
https://arxiv.org/abs/2310.06825
Dan Villablanca

FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness
Dao et al.
https://arxiv.org/abs/2205.14135
Benjamín Prieto

****
****
****
Semana Receso
****
****
****

****
Lecturas 23/09 Tokenization and Scaling Laws
****
Scaling Laws with Vocabulary: Larger Models Deserve Larger Vocabularies
Tao et al.
https://arxiv.org/abs/2407.13623

Training Compute-Optimal Large Language Models
Jordan Hoffmann et al.
https://arxiv.org/abs/2203.15556
Luz Infante

Stop Taking Tokenizers for Granted: They Are Core Design Decisions in Large Language Models
Alqahtani et al.
https://aclanthology.org/2026.eacl-long.394/
Valentina Juri

****
Lecturas 28/09 CPT
****
LatamGPT 70B Technical Report
Flores et al.
https://www.latamgpt.org/content/latamgpt-report.pdf

ALLaM: Large Language Models for Arabic and English
Saiful Bari et al.
https://arxiv.org/abs/2407.15390
Pablo Torres

Reuse, Don't Retrain: A Recipe for Continued Pretraining of Language Models
J. Parmar, S. Satheesh, M. Patwary, M. Shoeybi, and B. Catanzaro
https://arxiv.org/abs/2407.07263

Gabriel Venegas

****
Lecturas 30/09 Finetunning and MoE
****
Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity
W. Fedus, B, Zoph, and N. Shazeer
https://arxiv.org/abs/2101.03961

LoRA: Low-Rank Adaptation of Large Language Models
Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, Weizhu Chen
https://arxiv.org/abs/2106.09685
Joaquín Fuentealba

DeepSeekMoE: Towards Ultimate Expert Specialization in Mixture-of-Experts Language Models
Dai et al.
https://arxiv.org/abs/2401.06066
Ethan Orellana

****
Lecturas 05/10 Constrastive
****
Learning Transferable Visual Models From Natural Language Supervision
Alec Radford et al.
https://arxiv.org/abs/2103.00020

A Simple Framework for Contrastive Learning of Visual Representations
Ting Chen, Simon Kornblith, Mohammad Norouzi, Geoffrey Hinton
https://arxiv.org/abs/2002.05709
Matías Retamales

Self-Supervised Learning from Images with a Joint-Embedding Predictive Architecture
Assran et al.
https://arxiv.org/abs/2301.08243
Juan Antonio Moya

****
Lecturas 07/10 RAG
****
Improving Language Models by Retrieving from Trillions of Tokens
Sebastian Borgeaud et al.
https://arxiv.org/abs/2112.04426

REPLUG : Retrieval-Augmented Black-Box Language Models
Weijia Shi, et al.
https://arxiv.org/abs/2301.12652
Arturo Herreros

Sufficient Context: A New Lens on Retrieval Augmented Generation Systems
Hailey Joren et al.
https://arxiv.org/abs/2411.06037
Sebastián Palma

****
Lecturas 14/10 Metalearning
****
Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks
Chelsea Finn, Pieter Abbeel, Sergey Levine
https://arxiv.org/pdf/1703.03400.pdf

Transformers Learn In-Context by Gradient Descent
Von Oswald et al.
https://arxiv.org/abs/2212.07677
Christopher Stears

Rapid Task-Solving in Novel Environments
Sam Ritter, Ryan Faulkner, Laurent Sartran, Adam Santoro, Matt Botvinick, David Raposo
https://arxiv.org/abs/2006.03662
Clemente Valenzuela

****
Lecturas 19/10 Hallucinations in LLMs
****
Why Language Models Hallucinate
A. Tauman Kalai, O. Nachum, S. S. Vempala, and E. Zhang
https://arxiv.org/abs/2509.04664

Detecting hallucinations in large language models using semantic entropy
S. Farquhar, J. Kossen, L. Kuhn, and Y. Gal
https://www.nature.com/articles/s41586-024-07421-0
Rodrigo Meza

Hallucination-Free? Assessing the Reliability of Leading AI Legal Research Tools
V. Magesh, F. Surani, M. Dahl, M. Suzgun, C. Manning, D. Ho
https://arxiv.org/abs/2405.20362
Lilly Guevara

****
Lecturas 21/10 Models internal knowledge
****
Estimating Knowledge in Large Language Models Without Generating a Single Token
D. Gottesman and M. Geva
https://arxiv.org/abs/2406.12673

A Structural Probe for Finding Syntax in Word Representations,
John Hewitt and Christopher D. Manning
https://aclanthology.org/N19-1419.pdf
Tomás Ariztía

Locating and Editing Factual Associations in GPT
K. Meng, D. Bau, A. Andonian, and Y. Belinkov
https://arxiv.org/abs/2202.05262
Cristian Hinostroza

****
****
****
Término Ciclo 1

****
****
****