****
****
****
Ciclo 2
****
****
****

****
Lecturas 26/10 RL for LLMs post-training
****
Grupal
Set slides: RL4LLMs.pdf disponible web del curso

DAPO: An Open-Source LLM Reinforcement Learning System at Scale
Q. Yu et al. (2025)
https://arxiv.org/abs/2503.14476
Antonio Araujo

Does RL Really Incentivize Reasoning Capacity Beyond the Base Model?
Yang Yue, Zhiqi Chen, Rui Lu, Andrew Zhao, Zhaokai Wang, Yang Yue, Shiji Song, Gao Huang (2025)
https://arxiv.org/abs/2504.13837
Antonio Manriquez

A Minimalist Approach to LLM Reasoning: from Rejection Sampling to Reinforce
Wei Xiong, Jiarui Yao, Yuhui Xu, Bo Pang, Lei Wang, Doyen Sahoo, Junnan Li, Nan Jiang, Tong Zhang, Caiming Xiong, Hanze Dong
https://arxiv.org/abs/2504.11343
Marcelo Vargas

****
Lecturas 28/10 Agent Gyms
****
Paper Grupal
Reasoning Gym
Z. Stojanovski, O.r Stanley, J. Sharratt, R. Jones, A. Adefioye, J. Kaddour and A. Köpf (2025)
https://arxiv.org/abs/2505.24760

R2E-Gym: Procedural Environments and Hybrid Verifiers for Scaling Open-Weights SWE Agents
Naman Jain, Jaskirat Singh, Manish Shetty, Tianjun Zhang, Liang Zheng, Koushik Sen, and Ion Stoica (COLM 2025)
https://arxiv.org/abs/2504.07164

Evaluating and Training Large Language Model-based Agents across Diverse Environments
Z. Xi et al. ACL 2025
https://aclanthology.org/2025.acl-long.1355/

Breaking Barriers — Do Reinforcement Post Training Gains Transfer To Unseen Domains?
Chuxuan Hu, Yuxuan Zhu, Antony Kellermann, Caleb Biddulph, Suppakit Waiwitlikhit, Jason Benn, Daniel Kang (ICLR 2026)
https://arxiv.org/abs/2506.19733

****
Lecturas 02/11 Harness

****
Paper Grupal
SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering
John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, Ofir Press (Neurips 2024)
https://arxiv.org/abs/2405.15793

Darwin Godel Machine: Open-Ended Evolution of Self-Improving Agents
Jenny Zhang, Shengran Hu, Cong Lu, Robert Lange, Jeff Clune (ICLR, 2026)
https://arxiv.org/abs/2505.22954v3

SWE-Bench+: Enhanced Coding Benchmark for LLMs
Reem Aleithan, Haoran Xue, Mohammad Mahdi Mohajer, Elijah Nnorom, Gias Uddin, Song Wang (2024)
https://arxiv.org/abs/2410.06992

****
Lecturas 04/11 Agents
****
Paper Grupal
ReAct: Synergizing Reasoning and Acting in Language Models
Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, Yuan Cao (ICLR,2023)
https://arxiv.org/abs/2210.03629

ReasoningBank: Scaling Agent Self-Evolving with Reasoning Memory
Siru Ouyang, Jun Yan, I-Hung Hsu, Yanfei Chen, Ke Jiang, Zifeng Wang, Rujun Han, Long T. Le, Samira Daruki, Xiangru Tang, Vishy Tirumalashetty, George Lee, Mahsan Rofouei, Hangfei Lin, Jiawei Han, Chen-Yu Lee, Tomas Pfister   (ICLR 2026)
https://arxiv.org/abs/2509.25140

τ-bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains
Shunyu Yao, Noah Shinn, Pedram Razavi, Karthik Narasimhan (2024)
https://arxiv.org/abs/2406.12045

****
Lecturas 09/11 Reward Hacking
****
Paper Grupal
Defining and Characterizing Reward Gaming
Joar Skalse, Nikolaus H. R. Howe, Dmitrii Krasheninnikov, David Krueger (NEURIPS, 2022)
https://arxiv.org/abs/2209.13085

Scaling Laws for Reward Model Overoptimization
Leo Gao, John Schulman, Jacob Hilton (ICML 2023)
https://arxiv.org/abs/2210.10760

Monitoring Reasoning Models for Misbehavior and the Risks of Promoting Obfuscation
B. Baker, J. Huizinga, L. Gao, Z. Dou, M. Y. Guan, A. Madry, W. Zaremba, J. Pachocki, D. Farhi (2025)
https://arxiv.org/abs/2503.11926

****
Lecturas 11/11 Generative space (Tree of thoughts)
****
Paper Grupal

Mastering the game of Go with deep neural networks and tree search
David Silver et al.
https://www.nature.com/articles/nature16961

Tree of Thoughts: Deliberate Problem Solving with Large Language Models
Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, Karthik Narasimhan
https://arxiv.org/abs/2305.10601

Reasoning in Token Economies: Budget-Aware Evaluation of LLM Reasoning Strategies
Junlin Wang, Siddhartha Jain, Dejiao Zhang, Baishakhi Ray, Varun Kumar, Ben Athiwaratkun (EMNLP, 2024)
https://arxiv.org/abs/2406.06461

****
Lecturas 16/11 Implementation -> Inference
****
Paper Grupal
Efficient Memory Management for Large Language Model Serving with PagedAttention
Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, Ion Stoica (SOSP 2023)
https://arxiv.org/abs/2309.06180

Fast Inference from Transformers via Speculative Decoding
Yaniv Leviathan, Matan Kalman, Yossi Matias (ICML 2023)
https://arxiv.org/abs/2211.17192

DistServe: Disaggregating Prefill and Decoding for Goodput-optimized Large Language Model Serving
Yinmin Zhong, Shengyu Liu, Junda Chen, Jianbo Hu, Yibo Zhu, Xuanzhe Liu, Xin Jin, Hao Zhang (OSDI 2024)
https://arxiv.org/abs/2401.09670

An Efficient Inference Framework for Early-exit Large Language Models
Ruijie Miao, Yihan Yan, Xinshuo Yao, Tong Yang (2024)
https://arxiv.org/abs/2407.20272

****
Lecturas 18/11 Abstractions, Operators
****
Progress measures for grokking via mechanistic interpretability
Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, Jacob Steinhardt (2023)
https://arxiv.org/abs/2301.05217

Data Distributional Properties Drive Emergent In-Context Learning in Transformers
Stephanie Chan, Adam Santoro, Andrew Lampinen, Jane Wang, Aaditya Singh, Pierre Richemond, Jay McClelland, Felix Hill (2022)
https://arxiv.org/abs/2205.05055

Faith and Fate: Limits of Transformers on Compositionality
Nouha Dziri et al. (2023)
https://arxiv.org/abs/2305.18654

ARC-AGI-2: A New Challenge for Frontier AI Reasoning Systems
Francois Chollet, Mike Knoop, Gregory Kamradt, Bryan Landers, Henry Pinkard (2025)
https://arxiv.org/abs/2505.11831

****
Lecturas 23/11 World Models
****
Paper Grupal
Web Agents with World Models: Learning and Leveraging Environment Dynamics in Web Navigation
Hyungjoo Chae, Namyoung Kim, Kai Ong, Minju Gwak, Gwanwoo Song, Jihoon Kim, Sunghwan Kim, Dongha Lee, Jinyoung Yeo (ICLR 2025)
https://arxiv.org/abs/2410.13232

LLM-JEPA: Large Language Models Meet Joint Embedding Predictive Architectures
Hai Huang, Yann LeCun, Randall Balestriero (2025)
https://arxiv.org/abs/2509.14252

V-JEPA 2: Self-Supervised Video Models Enable Understanding, Prediction and Planning
Mido Assran et al. (2025)
https://arxiv.org/abs/2506.09985

****
Lecturas 25/11 Robotics
****
Paper grupal
π0.7 : a Steerable Generalist Robotic Foundation Model with Emergent Capabilities
A. Bo et al. (2026)
https://arxiv.org/abs/2604.15483

Do As I Can, Not As I Say: Grounding Language in Robotic Affordances
Michael Ahn et al. (2022)
https://arxiv.org/abs/2204.01691

Yell At Your Robot: Improving On-the-Fly from Language Corrections
Lucy Xiaoyang Shi, Zheyuan Hu, Tony Z. Zhao, Archit Sharma, Karl Pertsch, Jianlan Luo, Sergey Levine, Chelsea Finn (2024)
https://arxiv.org/abs/2403.12910

OpenVLA: An Open-Source Vision-Language-Action Model
Moo Jin Kim et al. (2024)
https://arxiv.org/abs/2406.09246